

## Problem Statement

You are a data scientist working for a school

You are asked to predict the GPA of the current students based on the following provided data:

0 StudentID int64  
1 Age int64  
2 Gender int64  
3 Ethnicity int64  
4 ParentalEducation int64  
5 StudyTimeWeekly float64 6 Absences int64  
7 Tutoring int64  
8 ParentalSupport int64  
9 Extracurricular int64  
10 Sports int64  
11 Music int64  
12 Volunteering int64  
13 GPA float64 14 GradeClass float64

The GPA is the Grade Point Average, typically ranges from 0.0 to 4.0 in most educational systems, with 4.0 representing an 'A' or excellent performance.

The minimum passing GPA can vary by institution, but it's often around 2.0. This usually corresponds to a 'C' grade, which is considered satisfactory.

You need to create a different Neural Network Architectures and compare your results. Predict the GPA of a Student based on a set of provided features. The data provided represents 2,392 students.

### 1) Import Libraries

In [175...

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Conv1D, MaxPooling1D, Flatten
from tensorflow.keras.regularizers import l2
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from tensorflow.keras.layers import BatchNormalization
```

### 2) Load Data

In [176...

```
# Load data
data = pd.read_csv("C:/Users/PC/OneDrive - Instituto Tecnológico y de Estudios S
data
```

Out[176...

	StudentID	Age	Gender	Ethnicity	ParentalEducation	StudyTimeWeekly	Absenc
0	1001	17	1	0	2	19.833723	
1	1002	18	0	0	1	15.408756	
2	1003	15	0	2	3	4.210570	
3	1004	17	1	0	3	10.028829	
4	1005	17	1	0	2	4.672495	
...	...	...	...	...	...	...	...
2387	3388	18	1	0	3	10.680555	
2388	3389	17	0	0	1	7.583217	
2389	3390	16	1	0	2	6.805500	
2390	3391	16	1	1	0	12.416653	
2391	3392	16	1	0	2	17.819907	

2392 rows × 15 columns



### 3) Review data:

In [177...

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2392 entries, 0 to 2391
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   StudentID             2392 non-null   int64
1   Age                   2392 non-null   int64
2   Gender                 2392 non-null   int64
3   Ethnicity              2392 non-null   int64
4   ParentalEducation      2392 non-null   int64
5   StudyTimeWeekly        2392 non-null   float64
6   Absences               2392 non-null   int64
7   Tutoring               2392 non-null   int64
8   ParentalSupport        2392 non-null   int64
9   Extracurricular        2392 non-null   int64
10  Sports                 2392 non-null   int64
11  Music                  2392 non-null   int64
12  Volunteering           2392 non-null   int64
13  GPA                    2392 non-null   float64
14  GradeClass             2392 non-null   float64
dtypes: float64(3), int64(12)
memory usage: 280.4 KB
```

## 4. Remove the columns not needed for Student performance prediction

```
In [178... # # Drop irrelevant columns
dataset = data.drop(['StudentID', 'Gender', 'Ethnicity', 'GradeClass'], axis=1)
dataset.head()
```

```
Out[178...   Age  ParentalEducation  StudyTimeWeekly  Absences  Tutoring  ParentalSupport  Ext
```

	Age	ParentalEducation	StudyTimeWeekly	Absences	Tutoring	ParentalSupport	Ext
0	17	2	19.833723	7	1	2	
1	18	1	15.408756	0	0	1	
2	15	3	4.210570	26	0	2	
3	17	3	10.028829	14	0	3	
4	17	2	4.672495	17	1	3	

## 5. Check if the columns has any null values:

```
In [179... # Check for null values
print("Conteo de valores nulos")
dataset.isnull().sum()
```

Conteo de valores nulos

```
Out[179... Age      0
ParentalEducation  0
StudyTimeWeekly   0
Absences          0
Tutoring          0
ParentalSupport   0
Extracurricular   0
Sports            0
Music             0
Volunteering      0
GPA               0
dtype: int64
```

## 6. Prepare your data for training and for testing set:

```
In [180... # Prepare the data
X = dataset.drop(['GPA'], axis=1)
y = data['GPA']

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

# Standardize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
# Check dimensions
print("Shape of X_train:", X_train.shape)
```

Shape of X\_train: (1913, 10)

## 7. Define different Neural Network Architectures

### Experiment 1: A single Dense Hidden Layer

```
In [181... # Define the neural network architecture
model = Sequential([
    Dense(64, activation='relu', input_shape=(X_train.shape[1],)),
    Dense(1, activation='linear') # Linear activation for regression
])

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mean_absol

model.summary()
```

c:\Users\PC\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential\_26"

Layer (type)	Output Shape	Param #
dense_84 (Dense)	(None, 64)	704
dense_85 (Dense)	(None, 1)	65

Total params: 769 (3.00 KB)

Trainable params: 769 (3.00 KB)

Non-trainable params: 0 (0.00 B)

```
In [182... # Train the model
history = model.fit(
    X_train,
    y_train,
    epochs=50, # Number of iterations
    batch_size=10, # Size of each batch
    validation_split=0.2, # Use 20% of the training data for validation
    verbose=1 # Display progress during training
)
```

Epoch 1/50

153/153 ————— 2s 4ms/step - loss: 2.9067 - mean\_absolute\_error: 1.4062 - val\_loss: 0.2592 - val\_mean\_absolute\_error: 0.4232  
Epoch 2/50

153/153 ————— 0s 2ms/step - loss: 0.1873 - mean\_absolute\_error: 0.3434 - val\_loss: 0.0886 - val\_mean\_absolute\_error: 0.2417  
Epoch 3/50

153/153 ————— 0s 2ms/step - loss: 0.0872 - mean\_absolute\_error: 0.2343 - val\_loss: 0.0713 - val\_mean\_absolute\_error: 0.2196  
Epoch 4/50

153/153 ————— 0s 2ms/step - loss: 0.0709 - mean\_absolute\_error: 0.2120 - val\_loss: 0.0609 - val\_mean\_absolute\_error: 0.2013  
Epoch 5/50

153/153 ————— 0s 2ms/step - loss: 0.0616 - mean\_absolute\_error: 0.2003 - val\_loss: 0.0537 - val\_mean\_absolute\_error: 0.1903  
Epoch 6/50

153/153 ————— 0s 2ms/step - loss: 0.0555 - mean\_absolute\_error: 0.1924 - val\_loss: 0.0495 - val\_mean\_absolute\_error: 0.1813  
Epoch 7/50

153/153 ————— 0s 2ms/step - loss: 0.0515 - mean\_absolute\_error: 0.1821 - val\_loss: 0.0479 - val\_mean\_absolute\_error: 0.1766  
Epoch 8/50

153/153 ————— 0s 2ms/step - loss: 0.0480 - mean\_absolute\_error: 0.1731 - val\_loss: 0.0469 - val\_mean\_absolute\_error: 0.1762  
Epoch 9/50

153/153 ————— 0s 2ms/step - loss: 0.0467 - mean\_absolute\_error: 0.1739 - val\_loss: 0.0458 - val\_mean\_absolute\_error: 0.1743  
Epoch 10/50

153/153 ————— 0s 2ms/step - loss: 0.0437 - mean\_absolute\_error: 0.1699 - val\_loss: 0.0449 - val\_mean\_absolute\_error: 0.1721  
Epoch 11/50

153/153 ————— 0s 2ms/step - loss: 0.0465 - mean\_absolute\_error: 0.1722 - val\_loss: 0.0428 - val\_mean\_absolute\_error: 0.1678  
Epoch 12/50

153/153 ————— 0s 2ms/step - loss: 0.0393 - mean\_absolute\_error: 0.1576 - val\_loss: 0.0431 - val\_mean\_absolute\_error: 0.1682  
Epoch 13/50

153/153 ————— 0s 2ms/step - loss: 0.0391 - mean\_absolute\_error: 0.1584 - val\_loss: 0.0433 - val\_mean\_absolute\_error: 0.1687  
Epoch 14/50

153/153 ————— 0s 2ms/step - loss: 0.0385 - mean\_absolute\_error: 0.1595 - val\_loss: 0.0424 - val\_mean\_absolute\_error: 0.1674  
Epoch 15/50

153/153 ————— 0s 2ms/step - loss: 0.0420 - mean\_absolute\_error: 0.1649 - val\_loss: 0.0425 - val\_mean\_absolute\_error: 0.1685  
Epoch 16/50

153/153 ————— 0s 2ms/step - loss: 0.0378 - mean\_absolute\_error: 0.1545 - val\_loss: 0.0442 - val\_mean\_absolute\_error: 0.1692  
Epoch 17/50

153/153 ————— 0s 2ms/step - loss: 0.0392 - mean\_absolute\_error: 0.1576 - val\_loss: 0.0409 - val\_mean\_absolute\_error: 0.1656  
Epoch 18/50

153/153 ————— 0s 2ms/step - loss: 0.0364 - mean\_absolute\_error: 0.1532 - val\_loss: 0.0420 - val\_mean\_absolute\_error: 0.1668  
Epoch 19/50

153/153 ————— 0s 2ms/step - loss: 0.0387 - mean\_absolute\_error: 0.1584 - val\_loss: 0.0456 - val\_mean\_absolute\_error: 0.1742  
Epoch 20/50

153/153 ————— 0s 2ms/step - loss: 0.0383 - mean\_absolute\_error: 0.1551 - val\_loss: 0.0418 - val\_mean\_absolute\_error: 0.1664  
Epoch 21/50

153/153 ————— 0s 2ms/step - loss: 0.0348 - mean\_absolute\_error: 0.1477 - val\_loss: 0.0439 - val\_mean\_absolute\_error: 0.1701  
Epoch 22/50

153/153 ————— 0s 3ms/step - loss: 0.0352 - mean\_absolute\_error: 0.1493 - val\_loss: 0.0423 - val\_mean\_absolute\_error: 0.1673  
Epoch 23/50

153/153 ————— 0s 2ms/step - loss: 0.0355 - mean\_absolute\_error: 0.1502 - val\_loss: 0.0439 - val\_mean\_absolute\_error: 0.1709  
Epoch 24/50

153/153 ————— 0s 2ms/step - loss: 0.0350 - mean\_absolute\_error: 0.1499 - val\_loss: 0.0413 - val\_mean\_absolute\_error: 0.1655  
Epoch 25/50

153/153 ————— 0s 2ms/step - loss: 0.0339 - mean\_absolute\_error: 0.1464 - val\_loss: 0.0432 - val\_mean\_absolute\_error: 0.1686  
Epoch 26/50

153/153 ————— 0s 2ms/step - loss: 0.0356 - mean\_absolute\_error: 0.1519 - val\_loss: 0.0425 - val\_mean\_absolute\_error: 0.1679  
Epoch 27/50

153/153 ————— 0s 2ms/step - loss: 0.0353 - mean\_absolute\_error: 0.1517 - val\_loss: 0.0437 - val\_mean\_absolute\_error: 0.1692  
Epoch 28/50

153/153 ————— 0s 2ms/step - loss: 0.0338 - mean\_absolute\_error: 0.1470 - val\_loss: 0.0428 - val\_mean\_absolute\_error: 0.1696  
Epoch 29/50

153/153 ————— 0s 2ms/step - loss: 0.0327 - mean\_absolute\_error: 0.1447 - val\_loss: 0.0422 - val\_mean\_absolute\_error: 0.1658  
Epoch 30/50

153/153 ————— 0s 2ms/step - loss: 0.0332 - mean\_absolute\_error: 0.1438 - val\_loss: 0.0426 - val\_mean\_absolute\_error: 0.1674  
Epoch 31/50

153/153 ————— 0s 2ms/step - loss: 0.0330 - mean\_absolute\_error: 0.1452 - val\_loss: 0.0425 - val\_mean\_absolute\_error: 0.1681  
Epoch 32/50

153/153 ————— 0s 2ms/step - loss: 0.0311 - mean\_absolute\_error: 0.1408 - val\_loss: 0.0418 - val\_mean\_absolute\_error: 0.1664  
Epoch 33/50

153/153 ————— 0s 2ms/step - loss: 0.0334 - mean\_absolute\_error: 0.1438 - val\_loss: 0.0426 - val\_mean\_absolute\_error: 0.1678  
Epoch 34/50

153/153 ————— 0s 2ms/step - loss: 0.0322 - mean\_absolute\_error: 0.1420 - val\_loss: 0.0486 - val\_mean\_absolute\_error: 0.1764  
Epoch 35/50

153/153 ————— 0s 2ms/step - loss: 0.0339 - mean\_absolute\_error: 0.1466 - val\_loss: 0.0430 - val\_mean\_absolute\_error: 0.1677  
Epoch 36/50

153/153 ————— 0s 2ms/step - loss: 0.0331 - mean\_absolute\_error: 0.1451 - val\_loss: 0.0427 - val\_mean\_absolute\_error: 0.1665  
Epoch 37/50

153/153 ————— 0s 2ms/step - loss: 0.0324 - mean\_absolute\_error: 0.1453 - val\_loss: 0.0438 - val\_mean\_absolute\_error: 0.1700  
Epoch 38/50

153/153 ————— 0s 2ms/step - loss: 0.0327 - mean\_absolute\_error: 0.1459 - val\_loss: 0.0416 - val\_mean\_absolute\_error: 0.1652  
Epoch 39/50

153/153 ————— 0s 2ms/step - loss: 0.0319 - mean\_absolute\_error: 0.1430 - val\_loss: 0.0441 - val\_mean\_absolute\_error: 0.1686  
Epoch 40/50

153/153 ————— 0s 2ms/step - loss: 0.0323 - mean\_absolute\_error: 0.1432 - val\_loss: 0.0422 - val\_mean\_absolute\_error: 0.1665  
Epoch 41/50

```

153/153 ————— 0s 2ms/step - loss: 0.0317 - mean_absolute_error: 0.
1427 - val_loss: 0.0425 - val_mean_absolute_error: 0.1666
Epoch 42/50
153/153 ————— 0s 2ms/step - loss: 0.0311 - mean_absolute_error: 0.
1388 - val_loss: 0.0427 - val_mean_absolute_error: 0.1682
Epoch 43/50
153/153 ————— 0s 2ms/step - loss: 0.0328 - mean_absolute_error: 0.
1454 - val_loss: 0.0424 - val_mean_absolute_error: 0.1677
Epoch 44/50
153/153 ————— 0s 2ms/step - loss: 0.0297 - mean_absolute_error: 0.
1373 - val_loss: 0.0425 - val_mean_absolute_error: 0.1668
Epoch 45/50
153/153 ————— 0s 2ms/step - loss: 0.0306 - mean_absolute_error: 0.
1387 - val_loss: 0.0430 - val_mean_absolute_error: 0.1696
Epoch 46/50
153/153 ————— 0s 2ms/step - loss: 0.0305 - mean_absolute_error: 0.
1410 - val_loss: 0.0433 - val_mean_absolute_error: 0.1691
Epoch 47/50
153/153 ————— 0s 2ms/step - loss: 0.0294 - mean_absolute_error: 0.
1362 - val_loss: 0.0418 - val_mean_absolute_error: 0.1667
Epoch 48/50
153/153 ————— 0s 2ms/step - loss: 0.0310 - mean_absolute_error: 0.
1426 - val_loss: 0.0447 - val_mean_absolute_error: 0.1720
Epoch 49/50
153/153 ————— 0s 2ms/step - loss: 0.0291 - mean_absolute_error: 0.
1370 - val_loss: 0.0456 - val_mean_absolute_error: 0.1735
Epoch 50/50
153/153 ————— 0s 2ms/step - loss: 0.0318 - mean_absolute_error: 0.
1408 - val_loss: 0.0450 - val_mean_absolute_error: 0.1712

```

In [183...

```

# History values for loss and MAE
loss = history.history['loss']
val_loss = history.history['val_loss']
mae = history.history['mean_absolute_error']
val_mae = history.history['val_mean_absolute_error']
epochs = range(1, len(loss) + 1)

# Training and Validation Loss Plot
plt.figure(figsize=(12, 5))

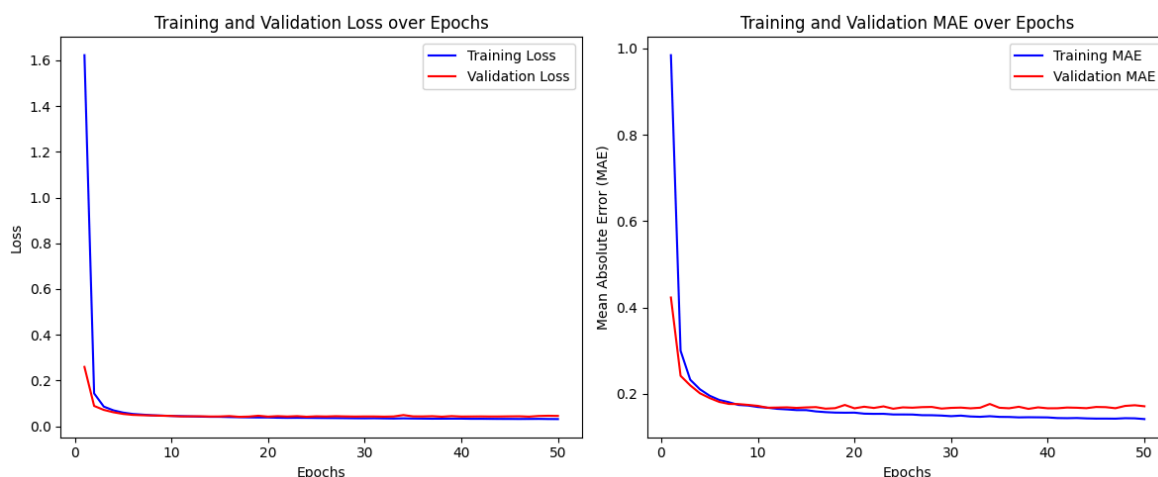
# Loss plot
plt.subplot(1, 2, 1)
plt.plot(epochs, loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'r', label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Validation Loss over Epochs')
plt.legend()

# MAE plot
plt.subplot(1, 2, 2)
plt.plot(epochs, mae, 'b', label='Training MAE')
plt.plot(epochs, val_mae, 'r', label='Validation MAE')
plt.xlabel('Epochs')
plt.ylabel('Mean Absolute Error (MAE)')
plt.title('Training and Validation MAE over Epochs')
plt.legend()

# Show plots

```

```
plt.tight_layout()
plt.show()
```



La gráfica refleja que el modelo del Experimento 1 con una sola capa densa es eficiente, con un aprendizaje rápido y estable. Además, demuestra una buena generalización a los datos de validación, ya que no hay signos de sobreajuste. Sin embargo, debido a la simplicidad del modelo, su capacidad para capturar patrones más complejos puede ser limitada en comparación con experimentos con arquitecturas más profundas o técnicas de regularización avanzadas.

In [184...

```
# Evaluate the model on the test set
test_loss, test_mae = model.evaluate(X_test, y_test, verbose=1)

print(f"Test Loss (MSE): {test_loss}")
print(f"Test Mean Absolute Error (MAE): {test_mae}")
```

15/15 ————— 0s 2ms/step - loss: 0.0506 - mean\_absolute\_error: 0.1777

Test Loss (MSE): 0.050356991589069366

Test Mean Absolute Error (MAE): 0.1774599552154541

In [185...

```
# Make predictions
y_pred = model.predict(X_test)

# Display the first 10 predictions and actual values for comparison
for i in range(10): # Display the first 10 predictions for comparison
    print(f"Predicted GPA: {y_pred[i][0]:.2f}, Actual GPA: {y_test.iloc[i]:.2f}")

# Calculate Overall MAE using sklearn
from sklearn.metrics import mean_absolute_error

# Calculate MAE between predictions and actual values
overall_mae = mean_absolute_error(y_test, y_pred)
print(f"\nOverall Mean Absolute Error (MAE) on Test Set: {overall_mae:.2f}")
```



15/15 ————— 0s 4ms/step

Predicted GPA: 1.63, Actual GPA: 1.43

Predicted GPA: 3.01, Actual GPA: 3.12

Predicted GPA: 1.81, Actual GPA: 2.04

Predicted GPA: 3.52, Actual GPA: 3.55

Predicted GPA: 0.14, Actual GPA: 0.25

Predicted GPA: 2.81, Actual GPA: 2.63

Predicted GPA: 1.65, Actual GPA: 2.06

Predicted GPA: 2.19, Actual GPA: 2.25

Predicted GPA: 2.08, Actual GPA: 2.19

Predicted GPA: 1.05, Actual GPA: 0.76

Overall Mean Absolute Error (MAE) on Test Set: 0.18

## Experiment 2: A set of three Dense Hidden Layers

In [186...

```
# Define the neural network architecture with three dense hidden layers
model_2 = Sequential([
    Dense(128, activation='relu', input_shape=(X_train.shape[1],)), # First hidden layer
    Dense(64, activation='relu'), # Second hidden layer
    Dense(32, activation='relu'), # Third hidden layer
    Dense(1, activation='linear') # Output layer for regression
])

# Compile the model
model_2.compile(optimizer='adam', loss='mean_squared_error', metrics=['mean_absolute_error'])

model_2.summary()
```

c:\Users\PC\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential\_27"

Layer (type)	Output Shape	Param #
dense_86 (Dense)	(None, 128)	1,408
dense_87 (Dense)	(None, 64)	8,256
dense_88 (Dense)	(None, 32)	2,080
dense_89 (Dense)	(None, 1)	33

Total params: 11,777 (46.00 KB)

Trainable params: 11,777 (46.00 KB)

Non-trainable params: 0 (0.00 B)

In [187...

```
# Train the model
history_2 = model_2.fit(
    X_train,
    y_train,
    epochs=50, # Number of iterations
```

```
batch_size=10,      # Size of each batch
validation_split=0.2, # Use 20% of the training data for validation
verbose=1           # Display progress during training
)
```

Epoch 1/50  
153/153 ————— 2s 4ms/step - loss: 0.8346 - mean\_absolute\_error: 0.6485 - val\_loss: 0.0791 - val\_mean\_absolute\_error: 0.2291  
Epoch 2/50  
153/153 ————— 0s 2ms/step - loss: 0.0692 - mean\_absolute\_error: 0.2124 - val\_loss: 0.0528 - val\_mean\_absolute\_error: 0.1890  
Epoch 3/50  
153/153 ————— 0s 2ms/step - loss: 0.0487 - mean\_absolute\_error: 0.1787 - val\_loss: 0.0566 - val\_mean\_absolute\_error: 0.1901  
Epoch 4/50  
153/153 ————— 0s 2ms/step - loss: 0.0461 - mean\_absolute\_error: 0.1722 - val\_loss: 0.0486 - val\_mean\_absolute\_error: 0.1759  
Epoch 5/50  
153/153 ————— 0s 2ms/step - loss: 0.0406 - mean\_absolute\_error: 0.1634 - val\_loss: 0.0581 - val\_mean\_absolute\_error: 0.1959  
Epoch 6/50  
153/153 ————— 0s 2ms/step - loss: 0.0400 - mean\_absolute\_error: 0.1601 - val\_loss: 0.0496 - val\_mean\_absolute\_error: 0.1791  
Epoch 7/50  
153/153 ————— 0s 2ms/step - loss: 0.0379 - mean\_absolute\_error: 0.1581 - val\_loss: 0.0574 - val\_mean\_absolute\_error: 0.1910  
Epoch 8/50  
153/153 ————— 0s 2ms/step - loss: 0.0386 - mean\_absolute\_error: 0.1562 - val\_loss: 0.0515 - val\_mean\_absolute\_error: 0.1781  
Epoch 9/50  
153/153 ————— 0s 2ms/step - loss: 0.0339 - mean\_absolute\_error: 0.1450 - val\_loss: 0.0524 - val\_mean\_absolute\_error: 0.1783  
Epoch 10/50  
153/153 ————— 0s 2ms/step - loss: 0.0321 - mean\_absolute\_error: 0.1439 - val\_loss: 0.0533 - val\_mean\_absolute\_error: 0.1828  
Epoch 11/50  
153/153 ————— 0s 2ms/step - loss: 0.0327 - mean\_absolute\_error: 0.1449 - val\_loss: 0.0508 - val\_mean\_absolute\_error: 0.1814  
Epoch 12/50  
153/153 ————— 0s 2ms/step - loss: 0.0328 - mean\_absolute\_error: 0.1445 - val\_loss: 0.0515 - val\_mean\_absolute\_error: 0.1796  
Epoch 13/50  
153/153 ————— 0s 2ms/step - loss: 0.0320 - mean\_absolute\_error: 0.1438 - val\_loss: 0.0501 - val\_mean\_absolute\_error: 0.1791  
Epoch 14/50  
153/153 ————— 0s 2ms/step - loss: 0.0316 - mean\_absolute\_error: 0.1419 - val\_loss: 0.0546 - val\_mean\_absolute\_error: 0.1849  
Epoch 15/50  
153/153 ————— 0s 2ms/step - loss: 0.0281 - mean\_absolute\_error: 0.1328 - val\_loss: 0.0515 - val\_mean\_absolute\_error: 0.1806  
Epoch 16/50  
153/153 ————— 0s 2ms/step - loss: 0.0284 - mean\_absolute\_error: 0.1336 - val\_loss: 0.0604 - val\_mean\_absolute\_error: 0.1988  
Epoch 17/50  
153/153 ————— 0s 2ms/step - loss: 0.0297 - mean\_absolute\_error: 0.1370 - val\_loss: 0.0544 - val\_mean\_absolute\_error: 0.1858  
Epoch 18/50  
153/153 ————— 0s 2ms/step - loss: 0.0307 - mean\_absolute\_error: 0.1388 - val\_loss: 0.0541 - val\_mean\_absolute\_error: 0.1824  
Epoch 19/50  
153/153 ————— 0s 2ms/step - loss: 0.0257 - mean\_absolute\_error: 0.1265 - val\_loss: 0.0511 - val\_mean\_absolute\_error: 0.1804  
Epoch 20/50  
153/153 ————— 0s 2ms/step - loss: 0.0264 - mean\_absolute\_error: 0.1281 - val\_loss: 0.0529 - val\_mean\_absolute\_error: 0.1807

Epoch 21/50  
153/153 ————— 0s 3ms/step - loss: 0.0276 - mean\_absolute\_error: 0.1319 - val\_loss: 0.0545 - val\_mean\_absolute\_error: 0.1837  
Epoch 22/50  
153/153 ————— 0s 2ms/step - loss: 0.0244 - mean\_absolute\_error: 0.1250 - val\_loss: 0.0564 - val\_mean\_absolute\_error: 0.1878  
Epoch 23/50  
153/153 ————— 0s 2ms/step - loss: 0.0248 - mean\_absolute\_error: 0.1244 - val\_loss: 0.0685 - val\_mean\_absolute\_error: 0.2073  
Epoch 24/50  
153/153 ————— 0s 2ms/step - loss: 0.0273 - mean\_absolute\_error: 0.1304 - val\_loss: 0.0531 - val\_mean\_absolute\_error: 0.1812  
Epoch 25/50  
153/153 ————— 0s 3ms/step - loss: 0.0225 - mean\_absolute\_error: 0.1181 - val\_loss: 0.0592 - val\_mean\_absolute\_error: 0.1948  
Epoch 26/50  
153/153 ————— 0s 3ms/step - loss: 0.0239 - mean\_absolute\_error: 0.1243 - val\_loss: 0.0498 - val\_mean\_absolute\_error: 0.1770  
Epoch 27/50  
153/153 ————— 0s 3ms/step - loss: 0.0227 - mean\_absolute\_error: 0.1195 - val\_loss: 0.0575 - val\_mean\_absolute\_error: 0.1876  
Epoch 28/50  
153/153 ————— 0s 2ms/step - loss: 0.0216 - mean\_absolute\_error: 0.1173 - val\_loss: 0.0575 - val\_mean\_absolute\_error: 0.1859  
Epoch 29/50  
153/153 ————— 0s 2ms/step - loss: 0.0226 - mean\_absolute\_error: 0.1194 - val\_loss: 0.0563 - val\_mean\_absolute\_error: 0.1875  
Epoch 30/50  
153/153 ————— 0s 3ms/step - loss: 0.0240 - mean\_absolute\_error: 0.1209 - val\_loss: 0.0604 - val\_mean\_absolute\_error: 0.1924  
Epoch 31/50  
153/153 ————— 0s 3ms/step - loss: 0.0221 - mean\_absolute\_error: 0.1164 - val\_loss: 0.0514 - val\_mean\_absolute\_error: 0.1814  
Epoch 32/50  
153/153 ————— 0s 3ms/step - loss: 0.0211 - mean\_absolute\_error: 0.1141 - val\_loss: 0.0560 - val\_mean\_absolute\_error: 0.1871  
Epoch 33/50  
153/153 ————— 0s 3ms/step - loss: 0.0198 - mean\_absolute\_error: 0.1113 - val\_loss: 0.0629 - val\_mean\_absolute\_error: 0.1995  
Epoch 34/50  
153/153 ————— 0s 3ms/step - loss: 0.0220 - mean\_absolute\_error: 0.1173 - val\_loss: 0.0569 - val\_mean\_absolute\_error: 0.1907  
Epoch 35/50  
153/153 ————— 0s 2ms/step - loss: 0.0201 - mean\_absolute\_error: 0.1121 - val\_loss: 0.0568 - val\_mean\_absolute\_error: 0.1894  
Epoch 36/50  
153/153 ————— 0s 3ms/step - loss: 0.0184 - mean\_absolute\_error: 0.1088 - val\_loss: 0.0611 - val\_mean\_absolute\_error: 0.1904  
Epoch 37/50  
153/153 ————— 1s 3ms/step - loss: 0.0210 - mean\_absolute\_error: 0.1135 - val\_loss: 0.0585 - val\_mean\_absolute\_error: 0.1879  
Epoch 38/50  
153/153 ————— 0s 3ms/step - loss: 0.0189 - mean\_absolute\_error: 0.1080 - val\_loss: 0.0619 - val\_mean\_absolute\_error: 0.1943  
Epoch 39/50  
153/153 ————— 0s 3ms/step - loss: 0.0205 - mean\_absolute\_error: 0.1129 - val\_loss: 0.0612 - val\_mean\_absolute\_error: 0.1942  
Epoch 40/50  
153/153 ————— 0s 3ms/step - loss: 0.0196 - mean\_absolute\_error: 0.1083 - val\_loss: 0.0596 - val\_mean\_absolute\_error: 0.1895

```

Epoch 41/50
153/153 ————— 0s 3ms/step - loss: 0.0164 - mean_absolute_error: 0.1003 - val_loss: 0.0563 - val_mean_absolute_error: 0.1879
Epoch 42/50
153/153 ————— 0s 3ms/step - loss: 0.0186 - mean_absolute_error: 0.1069 - val_loss: 0.0650 - val_mean_absolute_error: 0.1997
Epoch 43/50
153/153 ————— 0s 3ms/step - loss: 0.0178 - mean_absolute_error: 0.1048 - val_loss: 0.0662 - val_mean_absolute_error: 0.2027
Epoch 44/50
153/153 ————— 0s 3ms/step - loss: 0.0177 - mean_absolute_error: 0.1032 - val_loss: 0.0594 - val_mean_absolute_error: 0.1950
Epoch 45/50
153/153 ————— 0s 3ms/step - loss: 0.0199 - mean_absolute_error: 0.1095 - val_loss: 0.0618 - val_mean_absolute_error: 0.1967
Epoch 46/50
153/153 ————— 0s 3ms/step - loss: 0.0157 - mean_absolute_error: 0.0957 - val_loss: 0.0633 - val_mean_absolute_error: 0.1955
Epoch 47/50
153/153 ————— 0s 2ms/step - loss: 0.0161 - mean_absolute_error: 0.0990 - val_loss: 0.0717 - val_mean_absolute_error: 0.2098
Epoch 48/50
153/153 ————— 0s 2ms/step - loss: 0.0161 - mean_absolute_error: 0.0990 - val_loss: 0.0648 - val_mean_absolute_error: 0.1990
Epoch 49/50
153/153 ————— 0s 2ms/step - loss: 0.0162 - mean_absolute_error: 0.1011 - val_loss: 0.0642 - val_mean_absolute_error: 0.1975
Epoch 50/50
153/153 ————— 0s 2ms/step - loss: 0.0162 - mean_absolute_error: 0.0988 - val_loss: 0.0714 - val_mean_absolute_error: 0.2105

```

In [188...

```

# History values for Loss and MAE
loss_2 = history_2.history['loss']
val_loss_2 = history_2.history['val_loss']
mae_2 = history_2.history['mean_absolute_error']
val_mae_2 = history_2.history['val_mean_absolute_error']
epochs_2 = range(1, len(loss_2) + 1)

# Training and Validation Loss Plot
plt.figure(figsize=(12, 5))

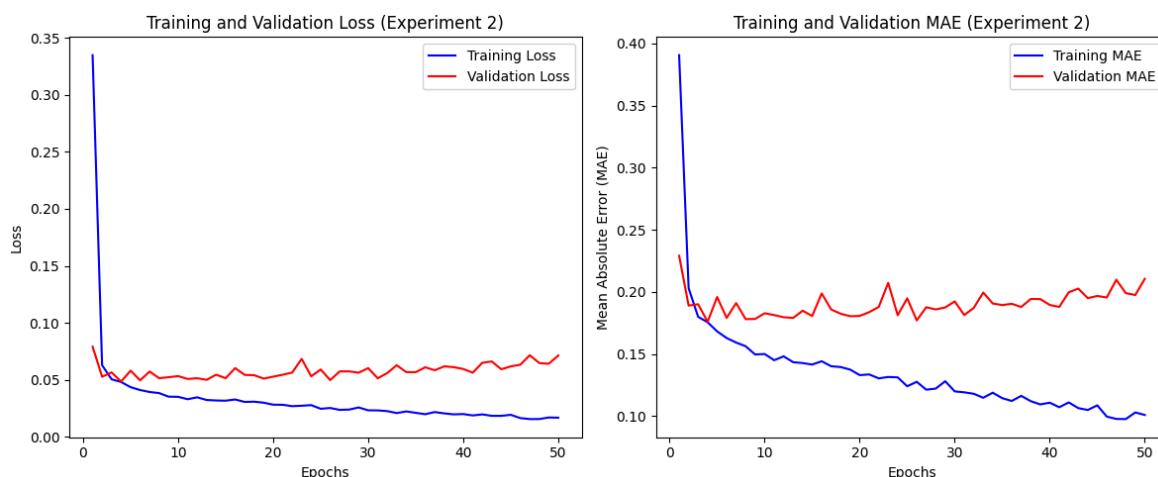
# Loss plot
plt.subplot(1, 2, 1)
plt.plot(epochs_2, loss_2, 'b', label='Training Loss')
plt.plot(epochs_2, val_loss_2, 'r', label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Validation Loss (Experiment 2)')
plt.legend()

# MAE plot
plt.subplot(1, 2, 2)
plt.plot(epochs_2, mae_2, 'b', label='Training MAE')
plt.plot(epochs_2, val_mae_2, 'r', label='Validation MAE')
plt.xlabel('Epochs')
plt.ylabel('Mean Absolute Error (MAE)')
plt.title('Training and Validation MAE (Experiment 2)')
plt.legend()

# Show plots

```

```
plt.tight_layout()
plt.show()
```



El modelo del Experimento 2 con tres capas densas tiene mayor capacidad de aprendizaje, pero la gráfica revela un problema de sobreajuste. Esto se evidencia por la creciente discrepancia entre las curvas de entrenamiento y validación. Para mejorar el desempeño de este modelo, sería necesario aplicar técnicas de regularización, como capas de Dropout o Batch Normalization, que podrían estabilizar el entrenamiento y mejorar la generalización.

In [189...

```
# Evaluate the model on the test set
test_loss_2, test_mae_2 = model_2.evaluate(X_test, y_test, verbose=1)
print(f"Experiment 2 - Test Loss (MSE): {test_loss_2}")
print(f"Experiment 2 - Test Mean Absolute Error (MAE): {test_mae_2}")

# Make predictions
y_pred_2 = model_2.predict(X_test)
```

```
15/15 ————— 0s 2ms/step - loss: 0.0736 - mean_absolute_error: 0.20
68
Experiment 2 - Test Loss (MSE): 0.07226863503456116
Experiment 2 - Test Mean Absolute Error (MAE): 0.20949308574199677
15/15 ————— 0s 10ms/step
```

In [190...

```
# Display the first 10 predictions and actual values for comparison
print(f"Experiment 2 - Predicted GPA: {y_pred_2[i][0]:.2f}, Actual GPA: {y_test.

# Calculate Overall MAE using sklearn
from sklearn.metrics import mean_absolute_error

# Calculate MAE between predictions and actual values
overall_mae_2 = mean_absolute_error(y_test, y_pred_2)
print(f"\nExperiment 2 - Overall Mean Absolute Error (MAE) on Test Set: {overall
```

```
Experiment 2 - Predicted GPA: 1.01, Actual GPA: 0.76
```

```
Experiment 2 - Overall Mean Absolute Error (MAE) on Test Set: 0.21
```

### Experiment 3: Add a dropout layer after each Dense Hidden Layer

In [191...

```
# Define the neural network architecture
model_3 = Sequential([
    Dense(128, activation='relu', input_shape=(X_train.shape[1],)), # First hid
```

```

Dropout(0.3), # Dropout layer to reduce overfitting
Dense(64, activation='relu'), # Second hidden layer
Dropout(0.3), # Dropout layer
Dense(32, activation='relu'), # Third hidden layer
Dropout(0.3), # Dropout layer
Dense(1, activation='linear') # Output layer for regression
])

# Compile the model
model_3.compile(optimizer='adam', loss='mean_squared_error', metrics=['mean_abs

model_3.summary()

```

c:\Users\PC\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential\_28"

Layer (type)	Output Shape	Param #
dense_90 (Dense)	(None, 128)	1,408
dropout_30 (Dropout)	(None, 128)	0
dense_91 (Dense)	(None, 64)	8,256
dropout_31 (Dropout)	(None, 64)	0
dense_92 (Dense)	(None, 32)	2,080
dropout_32 (Dropout)	(None, 32)	0
dense_93 (Dense)	(None, 1)	33

Total params: 11,777 (46.00 KB)

Trainable params: 11,777 (46.00 KB)

Non-trainable params: 0 (0.00 B)

In [192...

```

# Train the model
history_3 = model_3.fit(
    X_train,
    y_train,
    epochs=50,           # Number of iterations
    batch_size=10,       # Size of each batch
    validation_split=0.2, # Use 20% of the training data for validation
    verbose=1            # Display progress during training
)

```

Epoch 1/50

153/153 ————— 3s 4ms/step - loss: 2.3392 - mean\_absolute\_error: 1.2162 - val\_loss: 0.1580 - val\_mean\_absolute\_error: 0.3331  
Epoch 2/50

153/153 ————— 0s 3ms/step - loss: 0.4995 - mean\_absolute\_error: 0.5434 - val\_loss: 0.1702 - val\_mean\_absolute\_error: 0.3463  
Epoch 3/50

153/153 ————— 0s 3ms/step - loss: 0.3898 - mean\_absolute\_error: 0.4886 - val\_loss: 0.1210 - val\_mean\_absolute\_error: 0.2909  
Epoch 4/50

153/153 ————— 0s 3ms/step - loss: 0.3337 - mean\_absolute\_error: 0.4436 - val\_loss: 0.1009 - val\_mean\_absolute\_error: 0.2624  
Epoch 5/50

153/153 ————— 0s 3ms/step - loss: 0.2955 - mean\_absolute\_error: 0.4108 - val\_loss: 0.0983 - val\_mean\_absolute\_error: 0.2596  
Epoch 6/50

153/153 ————— 0s 3ms/step - loss: 0.2570 - mean\_absolute\_error: 0.3782 - val\_loss: 0.1103 - val\_mean\_absolute\_error: 0.2792  
Epoch 7/50

153/153 ————— 0s 3ms/step - loss: 0.2497 - mean\_absolute\_error: 0.3786 - val\_loss: 0.0691 - val\_mean\_absolute\_error: 0.2168  
Epoch 8/50

153/153 ————— 0s 3ms/step - loss: 0.2787 - mean\_absolute\_error: 0.4043 - val\_loss: 0.0773 - val\_mean\_absolute\_error: 0.2300  
Epoch 9/50

153/153 ————— 0s 3ms/step - loss: 0.2168 - mean\_absolute\_error: 0.3594 - val\_loss: 0.0611 - val\_mean\_absolute\_error: 0.2008  
Epoch 10/50

153/153 ————— 1s 4ms/step - loss: 0.2180 - mean\_absolute\_error: 0.3544 - val\_loss: 0.0863 - val\_mean\_absolute\_error: 0.2421  
Epoch 11/50

153/153 ————— 0s 3ms/step - loss: 0.2080 - mean\_absolute\_error: 0.3477 - val\_loss: 0.0890 - val\_mean\_absolute\_error: 0.2438  
Epoch 12/50

153/153 ————— 0s 3ms/step - loss: 0.1754 - mean\_absolute\_error: 0.3246 - val\_loss: 0.0864 - val\_mean\_absolute\_error: 0.2439  
Epoch 13/50

153/153 ————— 0s 3ms/step - loss: 0.1858 - mean\_absolute\_error: 0.3260 - val\_loss: 0.0611 - val\_mean\_absolute\_error: 0.2007  
Epoch 14/50

153/153 ————— 1s 3ms/step - loss: 0.1782 - mean\_absolute\_error: 0.3239 - val\_loss: 0.0890 - val\_mean\_absolute\_error: 0.2445  
Epoch 15/50

153/153 ————— 1s 3ms/step - loss: 0.1638 - mean\_absolute\_error: 0.3094 - val\_loss: 0.0579 - val\_mean\_absolute\_error: 0.1961  
Epoch 16/50

153/153 ————— 0s 3ms/step - loss: 0.1586 - mean\_absolute\_error: 0.3065 - val\_loss: 0.1137 - val\_mean\_absolute\_error: 0.2787  
Epoch 17/50

153/153 ————— 0s 3ms/step - loss: 0.1595 - mean\_absolute\_error: 0.3106 - val\_loss: 0.0537 - val\_mean\_absolute\_error: 0.1881  
Epoch 18/50

153/153 ————— 1s 3ms/step - loss: 0.1540 - mean\_absolute\_error: 0.3012 - val\_loss: 0.0507 - val\_mean\_absolute\_error: 0.1823  
Epoch 19/50

153/153 ————— 0s 3ms/step - loss: 0.1463 - mean\_absolute\_error: 0.2953 - val\_loss: 0.0523 - val\_mean\_absolute\_error: 0.1859  
Epoch 20/50

153/153 ————— 1s 3ms/step - loss: 0.1446 - mean\_absolute\_error: 0.2895 - val\_loss: 0.0489 - val\_mean\_absolute\_error: 0.1783  
Epoch 21/50



153/153 ————— 1s 3ms/step - loss: 0.1335 - mean\_absolute\_error: 0.2812 - val\_loss: 0.0909 - val\_mean\_absolute\_error: 0.2474  
Epoch 22/50

153/153 ————— 0s 3ms/step - loss: 0.1310 - mean\_absolute\_error: 0.2780 - val\_loss: 0.0616 - val\_mean\_absolute\_error: 0.2022  
Epoch 23/50

153/153 ————— 1s 3ms/step - loss: 0.1290 - mean\_absolute\_error: 0.2794 - val\_loss: 0.0519 - val\_mean\_absolute\_error: 0.1816  
Epoch 24/50

153/153 ————— 1s 3ms/step - loss: 0.1145 - mean\_absolute\_error: 0.2603 - val\_loss: 0.0565 - val\_mean\_absolute\_error: 0.1933  
Epoch 25/50

153/153 ————— 1s 3ms/step - loss: 0.1203 - mean\_absolute\_error: 0.2685 - val\_loss: 0.0577 - val\_mean\_absolute\_error: 0.1932  
Epoch 26/50

153/153 ————— 1s 3ms/step - loss: 0.1137 - mean\_absolute\_error: 0.2548 - val\_loss: 0.0797 - val\_mean\_absolute\_error: 0.2300  
Epoch 27/50

153/153 ————— 1s 3ms/step - loss: 0.1191 - mean\_absolute\_error: 0.2665 - val\_loss: 0.0652 - val\_mean\_absolute\_error: 0.2055  
Epoch 28/50

153/153 ————— 1s 3ms/step - loss: 0.1192 - mean\_absolute\_error: 0.2646 - val\_loss: 0.0624 - val\_mean\_absolute\_error: 0.2013  
Epoch 29/50

153/153 ————— 1s 3ms/step - loss: 0.1081 - mean\_absolute\_error: 0.2563 - val\_loss: 0.0583 - val\_mean\_absolute\_error: 0.1928  
Epoch 30/50

153/153 ————— 1s 3ms/step - loss: 0.1134 - mean\_absolute\_error: 0.2560 - val\_loss: 0.0733 - val\_mean\_absolute\_error: 0.2194  
Epoch 31/50

153/153 ————— 0s 3ms/step - loss: 0.1033 - mean\_absolute\_error: 0.2425 - val\_loss: 0.0619 - val\_mean\_absolute\_error: 0.2016  
Epoch 32/50

153/153 ————— 0s 3ms/step - loss: 0.0991 - mean\_absolute\_error: 0.2405 - val\_loss: 0.0646 - val\_mean\_absolute\_error: 0.2071  
Epoch 33/50

153/153 ————— 1s 3ms/step - loss: 0.1098 - mean\_absolute\_error: 0.2557 - val\_loss: 0.0691 - val\_mean\_absolute\_error: 0.2100  
Epoch 34/50

153/153 ————— 1s 3ms/step - loss: 0.1015 - mean\_absolute\_error: 0.2464 - val\_loss: 0.0559 - val\_mean\_absolute\_error: 0.1933  
Epoch 35/50

153/153 ————— 0s 3ms/step - loss: 0.0960 - mean\_absolute\_error: 0.2373 - val\_loss: 0.0581 - val\_mean\_absolute\_error: 0.1945  
Epoch 36/50

153/153 ————— 0s 3ms/step - loss: 0.1059 - mean\_absolute\_error: 0.2490 - val\_loss: 0.0597 - val\_mean\_absolute\_error: 0.1957  
Epoch 37/50

153/153 ————— 0s 3ms/step - loss: 0.1029 - mean\_absolute\_error: 0.2488 - val\_loss: 0.0746 - val\_mean\_absolute\_error: 0.2186  
Epoch 38/50

153/153 ————— 1s 3ms/step - loss: 0.0944 - mean\_absolute\_error: 0.2377 - val\_loss: 0.0661 - val\_mean\_absolute\_error: 0.2066  
Epoch 39/50

153/153 ————— 1s 3ms/step - loss: 0.1058 - mean\_absolute\_error: 0.2532 - val\_loss: 0.0657 - val\_mean\_absolute\_error: 0.2051  
Epoch 40/50

153/153 ————— 1s 4ms/step - loss: 0.0949 - mean\_absolute\_error: 0.2355 - val\_loss: 0.0576 - val\_mean\_absolute\_error: 0.1915  
Epoch 41/50

```

153/153 ————— 1s 3ms/step - loss: 0.0990 - mean_absolute_error: 0.
2445 - val_loss: 0.0510 - val_mean_absolute_error: 0.1801
Epoch 42/50
153/153 ————— 1s 3ms/step - loss: 0.1015 - mean_absolute_error: 0.
2412 - val_loss: 0.0771 - val_mean_absolute_error: 0.2213
Epoch 43/50
153/153 ————— 1s 3ms/step - loss: 0.0959 - mean_absolute_error: 0.
2375 - val_loss: 0.0999 - val_mean_absolute_error: 0.2581
Epoch 44/50
153/153 ————— 0s 3ms/step - loss: 0.0842 - mean_absolute_error: 0.
2256 - val_loss: 0.0676 - val_mean_absolute_error: 0.2047
Epoch 45/50
153/153 ————— 0s 3ms/step - loss: 0.0856 - mean_absolute_error: 0.
2268 - val_loss: 0.0766 - val_mean_absolute_error: 0.2237
Epoch 46/50
153/153 ————— 0s 3ms/step - loss: 0.0966 - mean_absolute_error: 0.
2385 - val_loss: 0.0612 - val_mean_absolute_error: 0.1983
Epoch 47/50
153/153 ————— 1s 3ms/step - loss: 0.0857 - mean_absolute_error: 0.
2297 - val_loss: 0.0753 - val_mean_absolute_error: 0.2139
Epoch 48/50
153/153 ————— 1s 3ms/step - loss: 0.0809 - mean_absolute_error: 0.
2195 - val_loss: 0.0763 - val_mean_absolute_error: 0.2176
Epoch 49/50
153/153 ————— 0s 3ms/step - loss: 0.0901 - mean_absolute_error: 0.
2365 - val_loss: 0.0812 - val_mean_absolute_error: 0.2247
Epoch 50/50
153/153 ————— 0s 3ms/step - loss: 0.0937 - mean_absolute_error: 0.
2389 - val_loss: 0.1014 - val_mean_absolute_error: 0.2584

```

In [204...

```

# History values for loss and MAE
loss_3 = history_3.history['loss']
val_loss_3 = history_3.history['val_loss']
mae_3 = history_3.history['mean_absolute_error']
val_mae_3 = history_3.history['val_mean_absolute_error']
epochs_3 = range(1, len(loss_3) + 1)

# Training and Validation Loss Plot
plt.figure(figsize=(12, 5))

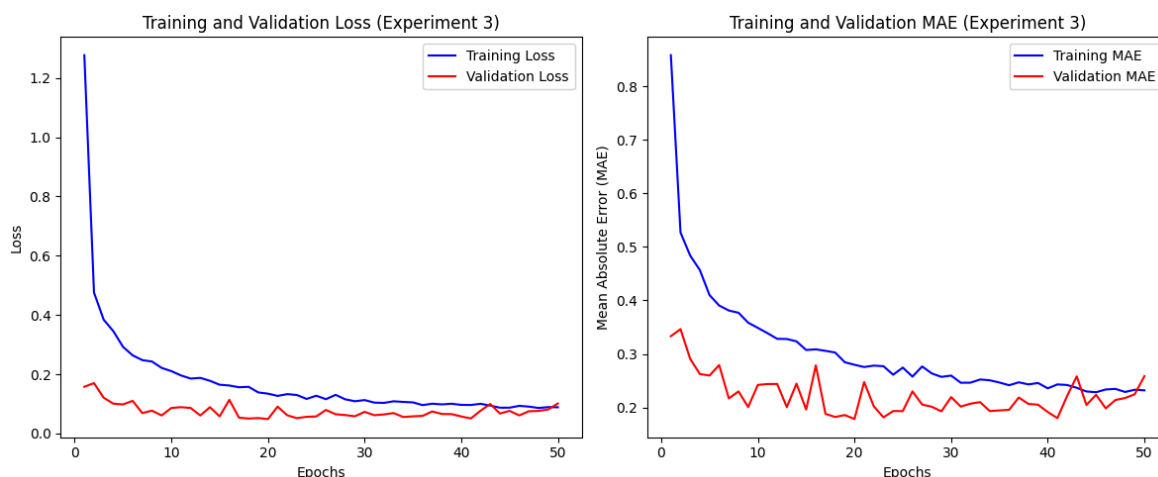
# Loss plot
plt.subplot(1, 2, 1)
plt.plot(epochs_3, loss_3, 'b', label='Training Loss')
plt.plot(epochs_3, val_loss_3, 'r', label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Validation Loss (Experiment 3)')
plt.legend()

# MAE plot
plt.subplot(1, 2, 2)
plt.plot(epochs_3, mae_3, 'b', label='Training MAE')
plt.plot(epochs_3, val_mae_3, 'r', label='Validation MAE')
plt.xlabel('Epochs')
plt.ylabel('Mean Absolute Error (MAE)')
plt.title('Training and Validation MAE (Experiment 3)')
plt.legend()

# Show plots

```

```
plt.tight_layout()
plt.show()
```



El modelo del Experimento 3, con capas de Dropout, muestra un desempeño más equilibrado que el modelo del Experimento 2, mitigando el sobreajuste y mejorando la generalización. Sin embargo, la introducción de Dropout parece haber reducido ligeramente la precisión del modelo en comparación con el Experimento 1, probablemente debido a que la regularización excesiva limita su capacidad de aprendizaje. Para optimizar este modelo, podría beneficiarse de combinaciones con técnicas adicionales, como normalización por lotes (Batch Normalization).

In [194...

```
# Evaluate the model on the test set
test_loss_3, test_mae_3 = model_3.evaluate(X_test, y_test, verbose=1)
print(f"Experiment 3 - Test Loss (MSE): {test_loss_3}")
print(f"Experiment 3 - Test Mean Absolute Error (MAE): {test_mae_3}")

# Make predictions
y_pred_3 = model_3.predict(X_test)
```

```
15/15 ————— 0s 2ms/step - loss: 0.0929 - mean_absolute_error: 0.24
85
Experiment 3 - Test Loss (MSE): 0.09313993155956268
Experiment 3 - Test Mean Absolute Error (MAE): 0.24728716909885406
15/15 ————— 0s 8ms/step
```

In [195...

```
# Display the first 10 predictions and actual values for comparison
print(f"Experiment 3 - Predicted GPA: {y_pred_3[i][0]:.2f}, Actual GPA: {y_test.

# Calculate Overall MAE using sklearn
from sklearn.metrics import mean_absolute_error

# Calculate MAE between predictions and actual values
overall_mae_3 = mean_absolute_error(y_test, y_pred_3)
print(f"\nExperiment 3 - Overall Mean Absolute Error (MAE) on Test Set: {overall
```

```
Experiment 3 - Predicted GPA: 1.12, Actual GPA: 0.76
```

```
Experiment 3 - Overall Mean Absolute Error (MAE) on Test Set: 0.25
```

**Experiment 4: Add a Batch Normalization Layer after each Dropout Layer.**

In [196...

```
# Define the neural network architecture for Experiment 4
model_4 = Sequential([
    Dense(128, activation='relu', input_shape=(X_train.shape[1],)), # First hidden layer
    Dropout(0.3), # Dropout layer
    BatchNormalization(), # Batch Normalization layer
    Dense(64, activation='relu'), # Second hidden layer
    Dropout(0.3), # Dropout layer
    BatchNormalization(), # Batch Normalization layer
    Dense(32, activation='relu'), # Third hidden layer
    Dropout(0.3), # Dropout layer
    BatchNormalization(), # Batch Normalization layer
    Dense(1, activation='linear') # Output layer for regression
])

# Compile the model
model_4.compile(optimizer='adam', loss='mean_squared_error', metrics=['mean_absolute_error'])

model_4.summary()
```

c:\Users\PC\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential\_29"

Layer (type)	Output Shape	Param #
dense_94 (Dense)	(None, 128)	1,408
dropout_33 (Dropout)	(None, 128)	0
batch_normalization_15 (BatchNormalization)	(None, 128)	512
dense_95 (Dense)	(None, 64)	8,256
dropout_34 (Dropout)	(None, 64)	0
batch_normalization_16 (BatchNormalization)	(None, 64)	256
dense_96 (Dense)	(None, 32)	2,080
dropout_35 (Dropout)	(None, 32)	0
batch_normalization_17 (BatchNormalization)	(None, 32)	128
dense_97 (Dense)	(None, 1)	33

Total params: 12,673 (49.50 KB)

Trainable params: 12,225 (47.75 KB)

Non-trainable params: 448 (1.75 KB)

In [197...

```
# Train the model
history_4 = model_4.fit(
    X_train,
    y_train,
    epochs=50,           # Number of iterations
    batch_size=10,       # Size of each batch
    validation_split=0.2, # Use 20% of the training data for validation
    verbose=1            # Display progress during training
)
```

Epoch 1/50  
153/153 ————— 6s 7ms/step - loss: 5.0640 - mean\_absolute\_error: 1.8581 - val\_loss: 0.9676 - val\_mean\_absolute\_error: 0.8109  
Epoch 2/50  
153/153 ————— 1s 4ms/step - loss: 1.7219 - mean\_absolute\_error: 1.0538 - val\_loss: 0.2404 - val\_mean\_absolute\_error: 0.4016  
Epoch 3/50  
153/153 ————— 1s 4ms/step - loss: 1.2176 - mean\_absolute\_error: 0.9033 - val\_loss: 0.1812 - val\_mean\_absolute\_error: 0.3415  
Epoch 4/50  
153/153 ————— 1s 4ms/step - loss: 0.9572 - mean\_absolute\_error: 0.7928 - val\_loss: 0.1250 - val\_mean\_absolute\_error: 0.2810  
Epoch 5/50  
153/153 ————— 1s 4ms/step - loss: 0.6811 - mean\_absolute\_error: 0.6579 - val\_loss: 0.1150 - val\_mean\_absolute\_error: 0.2727  
Epoch 6/50  
153/153 ————— 1s 4ms/step - loss: 0.5663 - mean\_absolute\_error: 0.6011 - val\_loss: 0.1023 - val\_mean\_absolute\_error: 0.2502  
Epoch 7/50  
153/153 ————— 1s 4ms/step - loss: 0.4338 - mean\_absolute\_error: 0.5179 - val\_loss: 0.0855 - val\_mean\_absolute\_error: 0.2316  
Epoch 8/50  
153/153 ————— 1s 4ms/step - loss: 0.4252 - mean\_absolute\_error: 0.5176 - val\_loss: 0.0775 - val\_mean\_absolute\_error: 0.2245  
Epoch 9/50  
153/153 ————— 1s 4ms/step - loss: 0.4041 - mean\_absolute\_error: 0.4968 - val\_loss: 0.0705 - val\_mean\_absolute\_error: 0.2146  
Epoch 10/50  
153/153 ————— 1s 4ms/step - loss: 0.3351 - mean\_absolute\_error: 0.4593 - val\_loss: 0.0665 - val\_mean\_absolute\_error: 0.2073  
Epoch 11/50  
153/153 ————— 1s 4ms/step - loss: 0.3275 - mean\_absolute\_error: 0.4554 - val\_loss: 0.0647 - val\_mean\_absolute\_error: 0.2047  
Epoch 12/50  
153/153 ————— 1s 4ms/step - loss: 0.3218 - mean\_absolute\_error: 0.4528 - val\_loss: 0.0619 - val\_mean\_absolute\_error: 0.2033  
Epoch 13/50  
153/153 ————— 1s 4ms/step - loss: 0.3255 - mean\_absolute\_error: 0.4494 - val\_loss: 0.0602 - val\_mean\_absolute\_error: 0.1960  
Epoch 14/50  
153/153 ————— 1s 4ms/step - loss: 0.2669 - mean\_absolute\_error: 0.4133 - val\_loss: 0.0621 - val\_mean\_absolute\_error: 0.2024  
Epoch 15/50  
153/153 ————— 1s 4ms/step - loss: 0.2805 - mean\_absolute\_error: 0.4156 - val\_loss: 0.0661 - val\_mean\_absolute\_error: 0.2107  
Epoch 16/50  
153/153 ————— 1s 4ms/step - loss: 0.2788 - mean\_absolute\_error: 0.4199 - val\_loss: 0.0553 - val\_mean\_absolute\_error: 0.1877  
Epoch 17/50  
153/153 ————— 1s 4ms/step - loss: 0.2695 - mean\_absolute\_error: 0.4141 - val\_loss: 0.0569 - val\_mean\_absolute\_error: 0.1940  
Epoch 18/50  
153/153 ————— 1s 4ms/step - loss: 0.2214 - mean\_absolute\_error: 0.3796 - val\_loss: 0.0526 - val\_mean\_absolute\_error: 0.1857  
Epoch 19/50  
153/153 ————— 1s 5ms/step - loss: 0.2460 - mean\_absolute\_error: 0.4007 - val\_loss: 0.0521 - val\_mean\_absolute\_error: 0.1882  
Epoch 20/50  
153/153 ————— 1s 4ms/step - loss: 0.2436 - mean\_absolute\_error: 0.3889 - val\_loss: 0.0482 - val\_mean\_absolute\_error: 0.1788

Epoch 21/50  
153/153 ————— 1s 4ms/step - loss: 0.2180 - mean\_absolute\_error: 0.3673 - val\_loss: 0.0522 - val\_mean\_absolute\_error: 0.1885  
Epoch 22/50  
153/153 ————— 1s 4ms/step - loss: 0.2382 - mean\_absolute\_error: 0.3988 - val\_loss: 0.0497 - val\_mean\_absolute\_error: 0.1796  
Epoch 23/50  
153/153 ————— 1s 4ms/step - loss: 0.2264 - mean\_absolute\_error: 0.3896 - val\_loss: 0.0514 - val\_mean\_absolute\_error: 0.1830  
Epoch 24/50  
153/153 ————— 1s 4ms/step - loss: 0.2282 - mean\_absolute\_error: 0.3789 - val\_loss: 0.0471 - val\_mean\_absolute\_error: 0.1789  
Epoch 25/50  
153/153 ————— 1s 4ms/step - loss: 0.2287 - mean\_absolute\_error: 0.3809 - val\_loss: 0.0470 - val\_mean\_absolute\_error: 0.1767  
Epoch 26/50  
153/153 ————— 1s 4ms/step - loss: 0.1969 - mean\_absolute\_error: 0.3545 - val\_loss: 0.0481 - val\_mean\_absolute\_error: 0.1775  
Epoch 27/50  
153/153 ————— 1s 4ms/step - loss: 0.2307 - mean\_absolute\_error: 0.3832 - val\_loss: 0.0497 - val\_mean\_absolute\_error: 0.1774  
Epoch 28/50  
153/153 ————— 1s 5ms/step - loss: 0.2166 - mean\_absolute\_error: 0.3668 - val\_loss: 0.0491 - val\_mean\_absolute\_error: 0.1807  
Epoch 29/50  
153/153 ————— 1s 4ms/step - loss: 0.2058 - mean\_absolute\_error: 0.3619 - val\_loss: 0.0485 - val\_mean\_absolute\_error: 0.1796  
Epoch 30/50  
153/153 ————— 1s 4ms/step - loss: 0.2072 - mean\_absolute\_error: 0.3594 - val\_loss: 0.0536 - val\_mean\_absolute\_error: 0.1915  
Epoch 31/50  
153/153 ————— 1s 4ms/step - loss: 0.2101 - mean\_absolute\_error: 0.3613 - val\_loss: 0.0512 - val\_mean\_absolute\_error: 0.1849  
Epoch 32/50  
153/153 ————— 1s 4ms/step - loss: 0.1987 - mean\_absolute\_error: 0.3505 - val\_loss: 0.0502 - val\_mean\_absolute\_error: 0.1812  
Epoch 33/50  
153/153 ————— 1s 5ms/step - loss: 0.1994 - mean\_absolute\_error: 0.3569 - val\_loss: 0.0499 - val\_mean\_absolute\_error: 0.1784  
Epoch 34/50  
153/153 ————— 1s 5ms/step - loss: 0.2149 - mean\_absolute\_error: 0.3710 - val\_loss: 0.0502 - val\_mean\_absolute\_error: 0.1780  
Epoch 35/50  
153/153 ————— 1s 5ms/step - loss: 0.2265 - mean\_absolute\_error: 0.3797 - val\_loss: 0.0525 - val\_mean\_absolute\_error: 0.1811  
Epoch 36/50  
153/153 ————— 1s 5ms/step - loss: 0.1872 - mean\_absolute\_error: 0.3449 - val\_loss: 0.0571 - val\_mean\_absolute\_error: 0.1885  
Epoch 37/50  
153/153 ————— 1s 5ms/step - loss: 0.2092 - mean\_absolute\_error: 0.3665 - val\_loss: 0.0550 - val\_mean\_absolute\_error: 0.1853  
Epoch 38/50  
153/153 ————— 1s 4ms/step - loss: 0.2147 - mean\_absolute\_error: 0.3717 - val\_loss: 0.0553 - val\_mean\_absolute\_error: 0.1934  
Epoch 39/50  
153/153 ————— 1s 4ms/step - loss: 0.1959 - mean\_absolute\_error: 0.3559 - val\_loss: 0.0532 - val\_mean\_absolute\_error: 0.1860  
Epoch 40/50  
153/153 ————— 1s 4ms/step - loss: 0.1974 - mean\_absolute\_error: 0.3602 - val\_loss: 0.0598 - val\_mean\_absolute\_error: 0.2003

Epoch 41/50  
**153/153** ————— 1s 4ms/step - loss: 0.2185 - mean\_absolute\_error: 0.3727 - val\_loss: 0.0543 - val\_mean\_absolute\_error: 0.1833  
 Epoch 42/50  
**153/153** ————— 1s 4ms/step - loss: 0.1923 - mean\_absolute\_error: 0.3491 - val\_loss: 0.0599 - val\_mean\_absolute\_error: 0.1911  
 Epoch 43/50  
**153/153** ————— 1s 4ms/step - loss: 0.2122 - mean\_absolute\_error: 0.3729 - val\_loss: 0.0532 - val\_mean\_absolute\_error: 0.1815  
 Epoch 44/50  
**153/153** ————— 1s 4ms/step - loss: 0.2004 - mean\_absolute\_error: 0.3572 - val\_loss: 0.0539 - val\_mean\_absolute\_error: 0.1838  
 Epoch 45/50  
**153/153** ————— 1s 4ms/step - loss: 0.1964 - mean\_absolute\_error: 0.3586 - val\_loss: 0.0536 - val\_mean\_absolute\_error: 0.1826  
 Epoch 46/50  
**153/153** ————— 1s 4ms/step - loss: 0.1795 - mean\_absolute\_error: 0.3392 - val\_loss: 0.0531 - val\_mean\_absolute\_error: 0.1845  
 Epoch 47/50  
**153/153** ————— 1s 4ms/step - loss: 0.1895 - mean\_absolute\_error: 0.3535 - val\_loss: 0.0669 - val\_mean\_absolute\_error: 0.2133  
 Epoch 48/50  
**153/153** ————— 1s 4ms/step - loss: 0.1866 - mean\_absolute\_error: 0.3422 - val\_loss: 0.0589 - val\_mean\_absolute\_error: 0.1974  
 Epoch 49/50  
**153/153** ————— 1s 4ms/step - loss: 0.2345 - mean\_absolute\_error: 0.3873 - val\_loss: 0.0578 - val\_mean\_absolute\_error: 0.1898  
 Epoch 50/50  
**153/153** ————— 1s 4ms/step - loss: 0.1890 - mean\_absolute\_error: 0.3507 - val\_loss: 0.0609 - val\_mean\_absolute\_error: 0.1920

In [206...

```
# History values for Loss and MAE
loss_4 = history_4.history['loss']
val_loss_4 = history_4.history['val_loss']
mae_4 = history_4.history['mean_absolute_error']
val_mae_4 = history_4.history['val_mean_absolute_error']
epochs_4 = range(1, len(loss_4) + 1)

# Training and Validation Loss Plot
plt.figure(figsize=(12, 5))

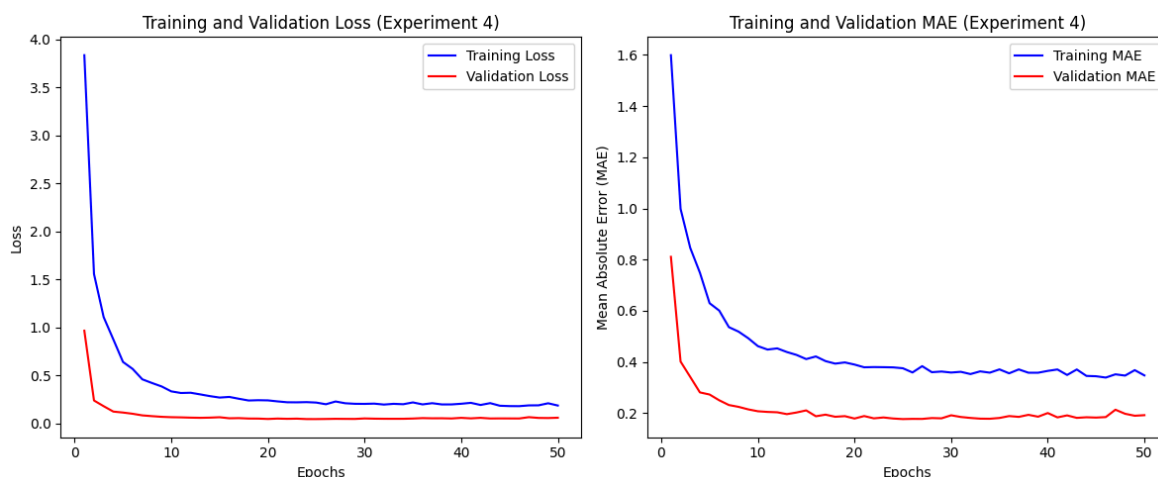
# Loss plot
plt.subplot(1, 2, 1)
plt.plot(epochs_4, loss_4, 'b', label='Training Loss')
plt.plot(epochs_4, val_loss_4, 'r', label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Validation Loss (Experiment 4)')
plt.legend()

# MAE plot
plt.subplot(1, 2, 2)
plt.plot(epochs_4, mae_4, 'b', label='Training MAE')
plt.plot(epochs_4, val_mae_4, 'r', label='Validation MAE')
plt.xlabel('Epochs')
plt.ylabel('Mean Absolute Error (MAE)')
plt.title('Training and Validation MAE (Experiment 4)')
plt.legend()

# Show plots
```



```
plt.tight_layout()
plt.show()
```



El Experimento 4, que incluye Dropout y Batch Normalization, muestra el mejor desempeño entre todos los experimentos. La gráfica indica un modelo bien regularizado, con curvas de pérdida y MAE muy cercanas entre entrenamiento y validación, lo que demuestra un excelente equilibrio entre aprendizaje y generalización. Esta combinación de técnicas permitió que el modelo aprovechara su capacidad de manera eficiente sin sobreajustarse.

In [199...

```
# Evaluate the model on the test set
test_loss_4, test_mae_4 = model_4.evaluate(X_test, y_test, verbose=1)
print(f"Experiment 4 - Test Loss (MSE): {test_loss_4}")
print(f"Experiment 4 - Test Mean Absolute Error (MAE): {test_mae_4}")

# Make predictions
y_pred_4 = model_4.predict(X_test)
```

```
15/15 ————— 0s 2ms/step - loss: 0.0650 - mean_absolute_error: 0.20
97
Experiment 4 - Test Loss (MSE): 0.058198168873786926
Experiment 4 - Test Mean Absolute Error (MAE): 0.19523905217647552
15/15 ————— 0s 12ms/step
```

In [200...

```
# Display the first 10 predictions and actual values for comparison
print(f"Experiment 4 - Predicted GPA: {y_pred_4[i][0]:.2f}, Actual GPA: {y_test.

# Calculate Overall MAE using sklearn
from sklearn.metrics import mean_absolute_error

# Calculate MAE between predictions and actual values
overall_mae_4 = mean_absolute_error(y_test, y_pred_4)
print(f"\nExperiment 4 - Overall Mean Absolute Error (MAE) on Test Set: {overall
```

```
Experiment 4 - Predicted GPA: 1.02, Actual GPA: 0.76
```

```
Experiment 4 - Overall Mean Absolute Error (MAE) on Test Set: 0.20
```

## Comparative table

In [207...

```
# Create a DataFrame focusing on the architecture of the neural networks
```

```
architecture_comparison = pd.DataFrame({
    "Experiment": [
        "Experiment 1: Single Dense Layer",
        "Experiment 2: Three Dense Layers",
        "Experiment 3: Dropout Layers",
        "Experiment 4: Dropout + Batch Norm"
    ],
    "Number of Hidden Layers": [1, 3, 3, 3],
    "Hidden Layers Configuration": [
        "1 Dense (64 neurons, ReLU)",
        "3 Dense (128, 64, 32 neurons, ReLU)",
        "3 Dense (128, 64, 32 neurons, ReLU) + Dropout (0.3)",
        "3 Dense (128, 64, 32 neurons, ReLU) + Dropout (0.3) + Batch Norm"
    ],
    "Regularization Methods": [
        "None",
        "None",
        "Dropout (0.3 after each hidden layer)",
        "Dropout (0.3) + Batch Normalization (after each hidden layer)"
    ],
    "Activation Function": ["ReLU", "ReLU", "ReLU", "ReLU"],
    "Output Layer": ["1 Dense (1 neuron, Linear)", "1 Dense (1 neuron, Linear)",
        "1 Dense (1 neuron, Linear)", "1 Dense (1 neuron, Linear)"]
})
```

architecture\_comparison

Out[207...

	Experiment	Number of Hidden Layers	Hidden Layers Configuration	Regularization Methods	Activation Function	Output Layer
0	Experiment 1: Single Dense Layer	1	1 Dense (64 neurons, ReLU)	None	ReLU	1 Dense (1 neuron, Linear)
1	Experiment 2: Three Dense Layers	3	3 Dense (128, 64, 32 neurons, ReLU)	None	ReLU	1 Dense (1 neuron, Linear)
2	Experiment 3: Dropout Layers	3	3 Dense (128, 64, 32 neurons, ReLU) + Dropout ...	Dropout (0.3 after each hidden layer)	ReLU	1 Dense (1 neuron, Linear)
3	Experiment 4: Dropout + Batch Norm	3	3 Dense (128, 64, 32 neurons, ReLU) + Dropout ...	Dropout (0.3) + Batch Normalization (after eac...	ReLU	1 Dense (1 neuron, Linear)

In [203...

```
# Create a DataFrame for the results of the experiments
experiment_results = pd.DataFrame({
    "Experiment": [
        "Experiment 1: Single Dense Layer",
        "Experiment 2: Three Dense Layers",
        "Experiment 3: Dropout Layers",
        "Experiment 4: Dropout + Batch Norm"
    ]
})
```

```

    ],
    "Hidden Layers Configuration": [
        "1 Dense (64 neurons, ReLU)",
        "3 Dense (128, 64, 32 neurons, ReLU)",
        "3 Dense (128, 64, 32 neurons, ReLU) + Dropout (0.3)",
        "3 Dense (128, 64, 32 neurons, ReLU) + Dropout (0.3) + Batch Norm"
    ],
    "Test Loss (MSE)": [
        test_loss,
        test_loss_2,
        test_loss_3,
        test_loss_4
    ],
    "Test MAE": [
        test_mae,
        test_mae_2,
        test_mae_3,
        test_mae_4
    ],
    "Overall MAE on Test Set": [
        overall_mae,
        overall_mae_2,
        overall_mae_3,
        overall_mae_4
    ]
}
))

experiment_results

```

Out[203...

	Experiment	Hidden Layers Configuration	Test Loss (MSE)	Test MAE	Overall MAE on Test Set
0	Experiment 1: Single Dense Layer	1 Dense (64 neurons, ReLU)	0.050357	0.177460	0.177460
1	Experiment 2: Three Dense Layers	3 Dense (128, 64, 32 neurons, ReLU)	0.072269	0.209493	0.209493
2	Experiment 3: Dropout Layers	3 Dense (128, 64, 32 neurons, ReLU) + Dropout ...	0.093140	0.247287	0.247287
3	Experiment 4: Dropout + Batch Norm	3 Dense (128, 64, 32 neurons, ReLU) + Dropout ...	0.058198	0.195239	0.195239

### 1. Experimento 1: Una Capa Densa Simple:

- Este experimento utilizó una arquitectura simple con una sola capa densa. Aunque obtuvo un desempeño razonable, la arquitectura podría carecer de la capacidad para aprender patrones complejos en los datos en comparación con modelos más profundos.

### 2. Experimento 2: Tres Capas Densas:

- Al aumentar el número de capas ocultas, el modelo incrementó su capacidad, pero obtuvo un rendimiento ligeramente peor que el Experimento 1. Esto

puede indicar sobreajuste o dificultades de optimización debido a la mayor complejidad del modelo.

### 3. Experimento 3: Capas de Dropout:

- La inclusión de capas de dropout mejoró la regularización del modelo, pero los valores más altos de pérdida y MAE sugieren un posible subajuste. El modelo podría haberse vuelto demasiado escaso debido al dropout, limitando su capacidad para capturar los patrones subyacentes.

### 4. Experimento 4: Dropout + Normalización por Lotes:

- La combinación de dropout con normalización por lotes resultó en el mejor desempeño de todos los experimentos. La normalización por lotes probablemente estabilizó y aceleró el entrenamiento, mientras que el dropout ayudó a evitar el sobreajuste, logrando un balance adecuado.

## Mejor Experimento:

El **Experimento 4: Dropout + Normalización por Lotes** es el mejor porque:

- Logró la **menor pérdida en pruebas (MSE)** (0.058198) y un **MAE general mejorado** (0.195239) en comparación con las demás arquitecturas.
- La combinación de normalización por lotes y dropout permitió regularizar el modelo eficazmente y estabilizar el aprendizaje, lo que resultó en una mejor capacidad de generalización.

## Conclusión:

Agregar **normalización por lotes** después de las capas de dropout (Experimento 4) ofreció la arquitectura más equilibrada, mitigando eficazmente el sobreajuste y manteniendo la capacidad del modelo para aprender relaciones complejas.